

Name: Shabbar
Roll No.: PB57
PRN: 1032221508

BDT ASSIGNMENT 4

HDFS Commands

```
*Commands X
[cloudera@quickstart ~]$ vi first.txt
[cloudera@quickstart ~]$ hdfs dfs -ls
[cloudera@quickstart ~]$ vi myfile.txt
[cloudera@quickstart ~]$ hdfs dfs -ls
[cloudera@quickstart ~]$ hdfs dfs -put first.txt
[cloudera@quickstart ~]$ hdfs dfs -ls
Found 1 items
-rw-r--r-- 1 cloudera cloudera      16 2024-08-20 22:50 first.txt
[cloudera@quickstart ~]$ hdfs dfs -Mkdir my_dir
-Mkdir: Unknown command
[cloudera@quickstart ~]$ hdfs dfs -mkdir my_dir
[cloudera@quickstart ~]$ hdfs dfs -put myfile.txt my_dir/
[cloudera@quickstart ~]$ hdfs dfs -ls
Found 2 items
-rw-r--r-- 1 cloudera cloudera      16 2024-08-20 22:50 first.txt
drwxr-xr-x - cloudera cloudera      0 2024-08-20 22:53 my_dir
[cloudera@quickstart ~]$ hadoop fs -mkdir dir
[cloudera@quickstart ~]$ hadoop fs -put sample.txt /user/cloudera
put: 'sample.txt': No such file or directory
[cloudera@quickstart ~]$ hadoop fs -put first.txt /user/cloudera
put: '/user/cloudera/first.txt': File exists
[cloudera@quickstart ~]$ hadoop fs -get /user/cloudera/first.txt desktop/
get: 'desktop/': No such file or directory
[cloudera@quickstart ~]$ hadoop fs -get first.txt
get: 'first.txt': No such file or directory
[cloudera@quickstart ~]$ hadoop fs -get first.txt
get: 'first.txt': File exists
[cloudera@quickstart ~]$ hdfs dfs -rm -R /user/cloudera/first.txt
24/08/20 23:10:33 INFO fs.TrashPolicyDefault: Namenode trash configuration: Deletion interval = 0 minutes, Emptier interval = 0 minutes.
Deleted /user/cloudera/first.txt
[cloudera@quickstart ~]$ hadoop fs -get first.txt
get: 'first.txt': No such file or directory
[cloudera@quickstart ~]$ hadoop fs -get my_dir/myfile.txt Desktop
[cloudera@quickstart ~]$ vi first.txt
[cloudera@quickstart ~]$ hdfs dfs -put first.txt
[cloudera@quickstart ~]$ hadoop fs -mv /user/cloudera/my_dir /user/cloudera/dir
[cloudera@quickstart ~]$
```



```

-s Rather than showing the size of each individual file that matches the
pattern, shows the total (summary) size.
-h Formats the sizes of files in a human-readable fashion rather than a number
of bytes.

Note that, even without the -s option, this only shows size summaries one level
deep into a directory.

The output is in the form
    size    disk space consumed    name(full path)

-expunge :
Empty the Trash

-get [-p] [-ignoreCrc] [-crc] <src> ... <localdst> :
Copy files that match the file pattern <src> to the local name. <src> is kept.
When copying multiple files, the destination must be a directory. Passing -p
preserves access and modification times, ownership and the mode.

-getfacl [-R] <path> :
Displays the Access Control Lists (ACLs) of files and directories. If a
directory has a default ACL, then getfacl also displays the default ACL.

-R List the ACLs of all files and directories recursively.
<path> File or directory to list.

-getfattr [-R] [-n name] [-d] [-e en] <path> :
Displays the extended attribute names and values (if any) for a file or
directory.

-R Recursively list the attributes for all files and directories.
-n name Dump the named extended attribute value.
-d Dump all extended attribute values associated with pathname.
-e <encoding> Encode values after retrieving them. Valid encodings are "text",
"hex", and "base64". Values encoded as text strings are enclosed
in double quotes ["], and values encoded as hexadecimal and
base64 are prefixed with 0x and 0s, respectively.
<path> The file or directory.

-getmerge [-n] <src> <localdst> :
Get all the files in the directories that match the source file pattern and
merge and sort them to only one file on local fs. <src> is kept.

-nl Add a newline character at the end of each file.

-help [cmd ...] :
Displays help for given command or all commands if none is specified.

-ls [-d] [-h] [-R] [<path> ...] :
List the contents that match the specified file pattern. If path is not
specified, the contents of /user/<currentUser> will be listed. Directory entries

```

```

If only the owner or group is specified, then only the owner or group is
modified. The owner and group names may only consist of digits, alphabet,
and any of [._-/@a-zA-Z0-9]. The names are case sensitive.

WARNING: Avoid using '.' to separate user name and group though Linux allows it.
If user names have dots in them and you are using local file system, you might
see surprising results since the shell command 'chown' is used for local files.

-copyFromLocal [-f] [-p] [-l] <localsrc> ... <dst> :
Identical to the -put command.

-copyToLocal [-p] [-ignoreCrc] [-crc] <src> ... <localdst> :
Identical to the -get command.

-count [-q] [-h] <paths> ... :
Count the number of directories, files and bytes under the paths
that match the specified file pattern. The output columns are:
DIR COUNT FILE COUNT CONTENT SIZE FILE NAME or
QUOTA REMAINING QUOTA SPACE QUOTA REMAINING SPACE QUOTA
DIR COUNT FILE COUNT CONTENT SIZE FILE NAME
The -h option shows file sizes in human readable format.

-cp [-f] [-p] [-p[topax]] <src> ... <dst> :
Copy files that match the file pattern <src> to a destination. When copying
multiple files, the destination must be a directory. Passing -p preserves status
[topax] (timestamps, ownership, permission, ACLs, XAttr). If -p is specified
with no <arg>, then preserves timestamps, ownership, permission. If -pa is
specified, then preserves permission also because ACL is a super-set of
permission. Passing -f overwrites the destination if it already exists. raw
namespace extended attributes are preserved if (1) they are supported (HDFS
only) and, (2) all of the source and target pathnames are in the /.reserved/raw
hierarchy. raw namespace xattr preservation is determined solely by the presence
(or absence) of the /.reserved/raw prefix and not by the -p option.

-createSnapshot <snapshotDir> [<snapshotName>] :
Create a snapshot on a directory

-deleteSnapshot <snapshotDir> <snapshotName> :
Delete a snapshot from a directory

-df [-h] [<path> ...] :
Shows the capacity, free and used space of the filesystem. If the filesystem has
multiple partitions, and no path to a particular partition is specified, then
the status of the root partitions will be shown.

-h Formats the sizes of files in a human-readable fashion rather than a number
of bytes.

-du [-s] [-h] [<path> ...] :
Show the amount of space, in bytes, used by the files that match the specified

```